Surviving the AI Slopageddon

How AI is shaking Open Source



Raphaël Semeteys

2026 October 7

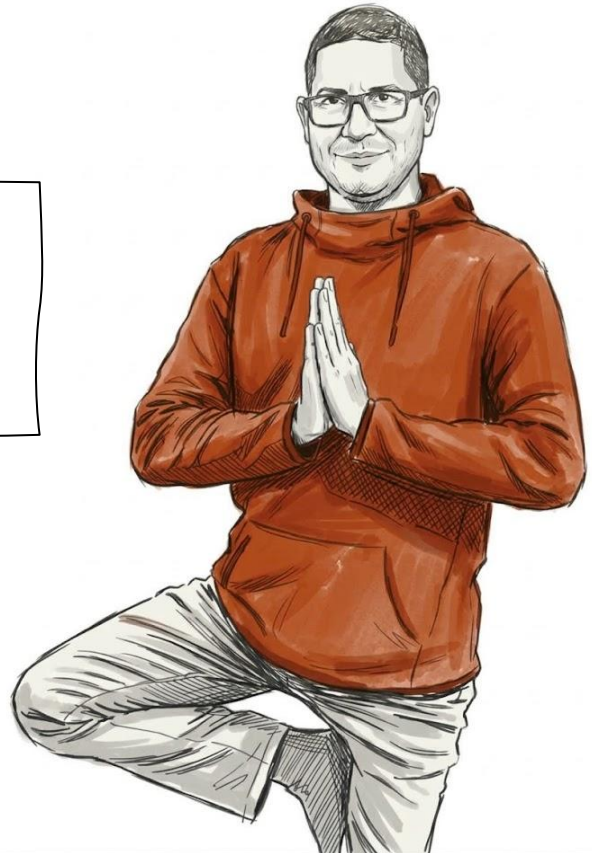
Raphaël Semeteys

WORLDLINE

Tech & AI Advocate
Open Source Expert

PROJETS LIBRES
LINUXFR.ORG

LinuxFr
Eclipse QSOS

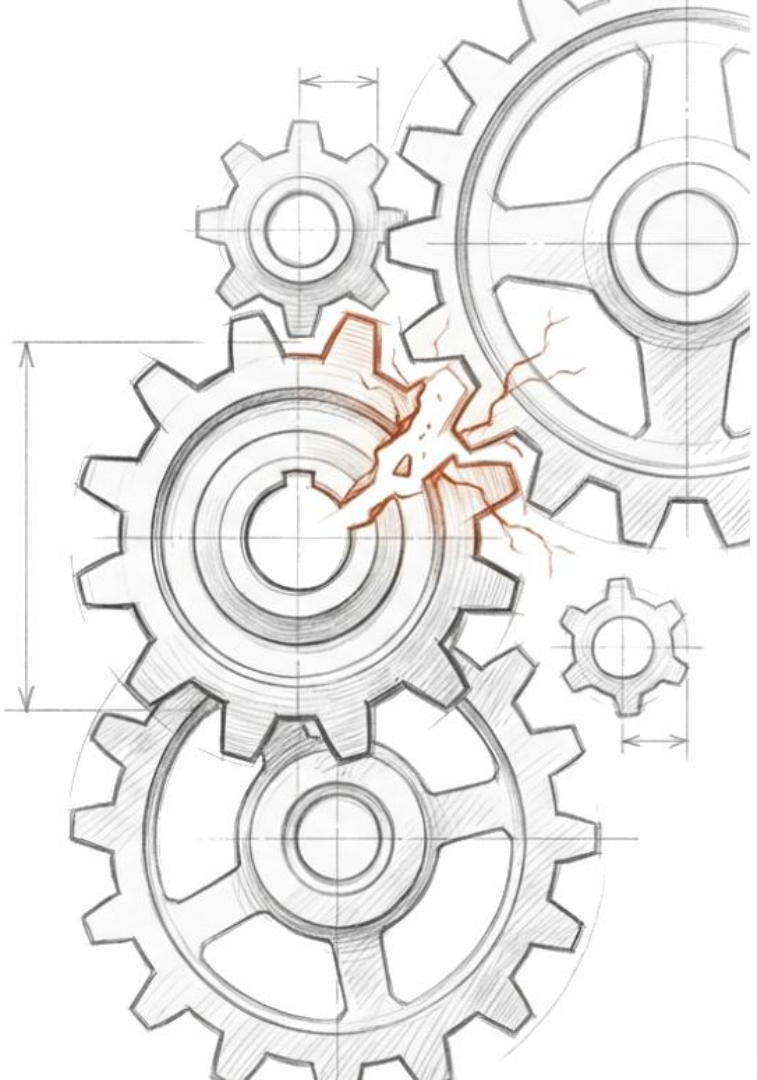


Also

Lore Architect (novels, games)

DEVFEST PARIS 2026

1. AI Shift
2. First symptoms
3. Is the system collapsing?
4. Adapting to an AI-Native world
5. What happens next?





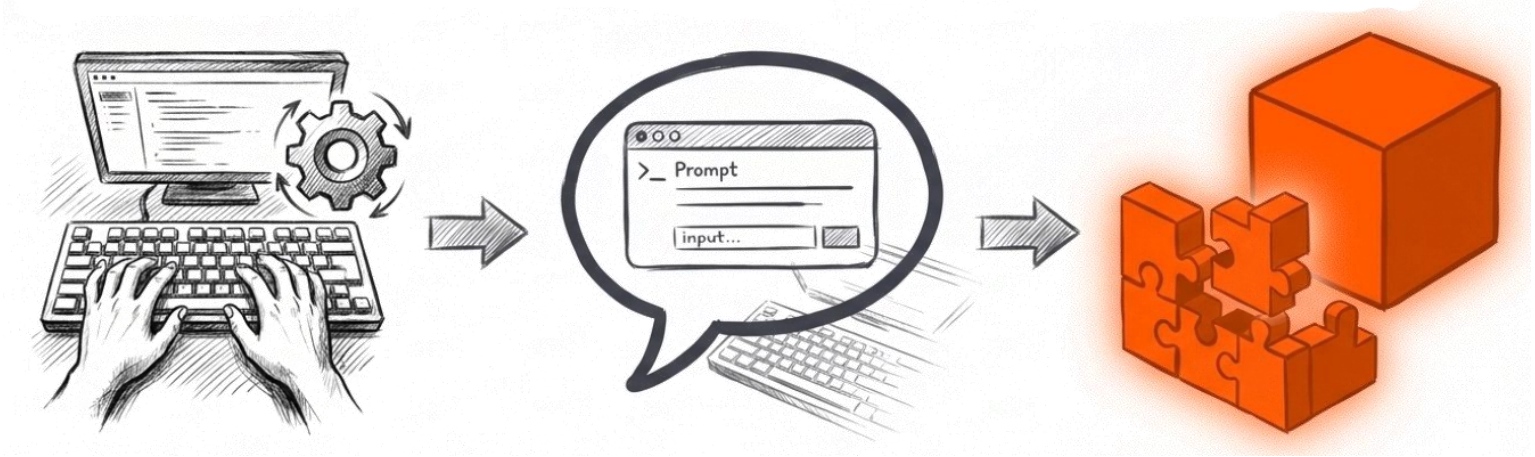
1. AI Shift

From code to prompts to agents

Assistance
Auto-completion

Natural Language
Mediation

Delegated
autonomous work



AI agents assemble open source software from prompts

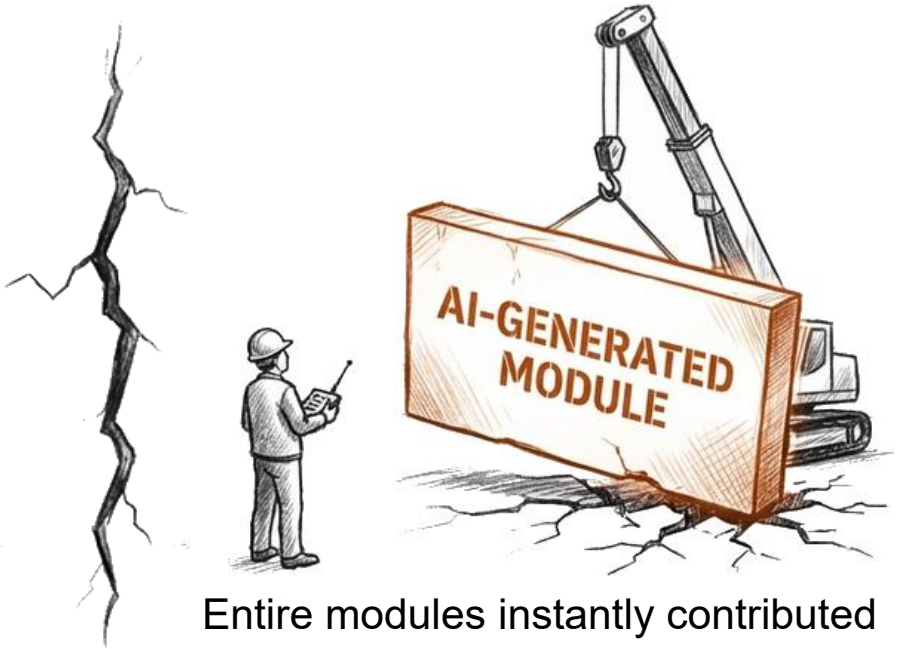
Open Source... without a community?

Collaboration brick by brick



Human-contributed patches
Building a community

Laying entire walls

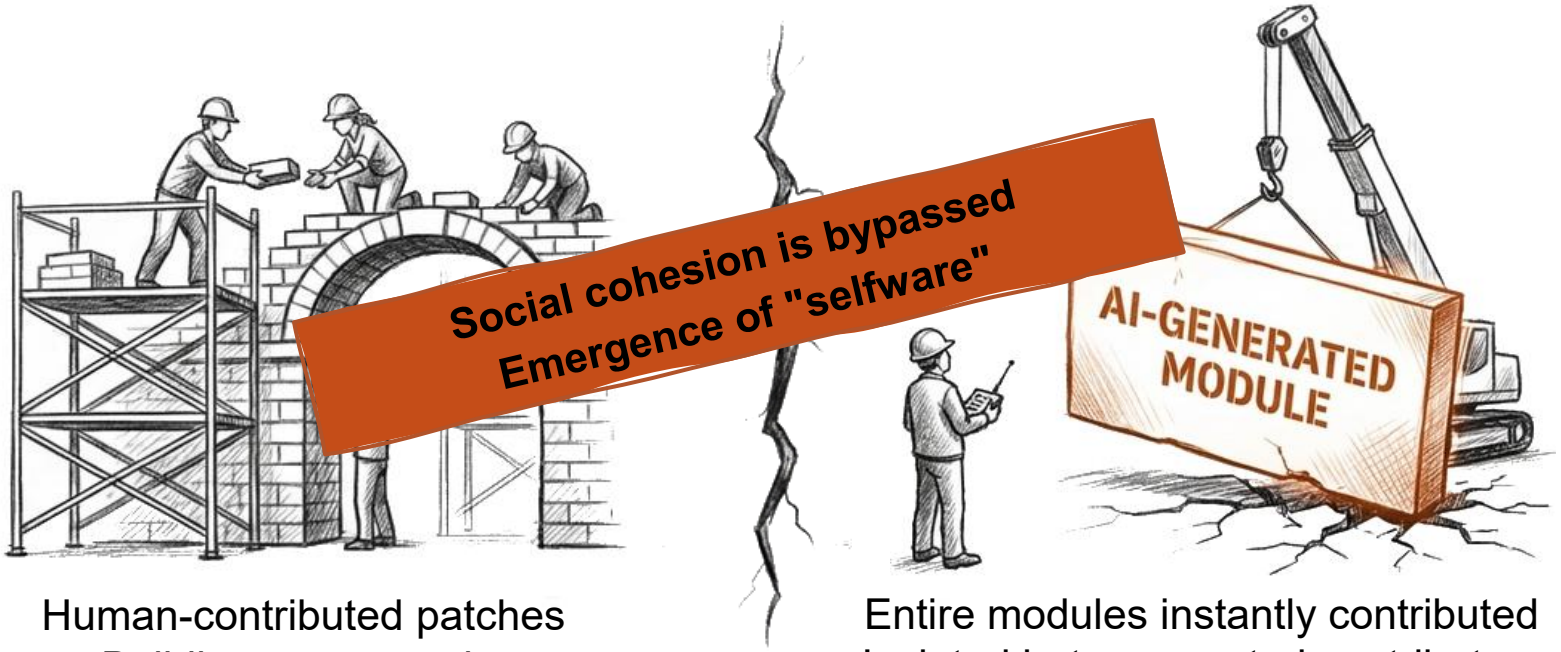


Entire modules instantly contributed
Isolated but augmented contributors

Open Source... without a community?

Collaboration brick by brick

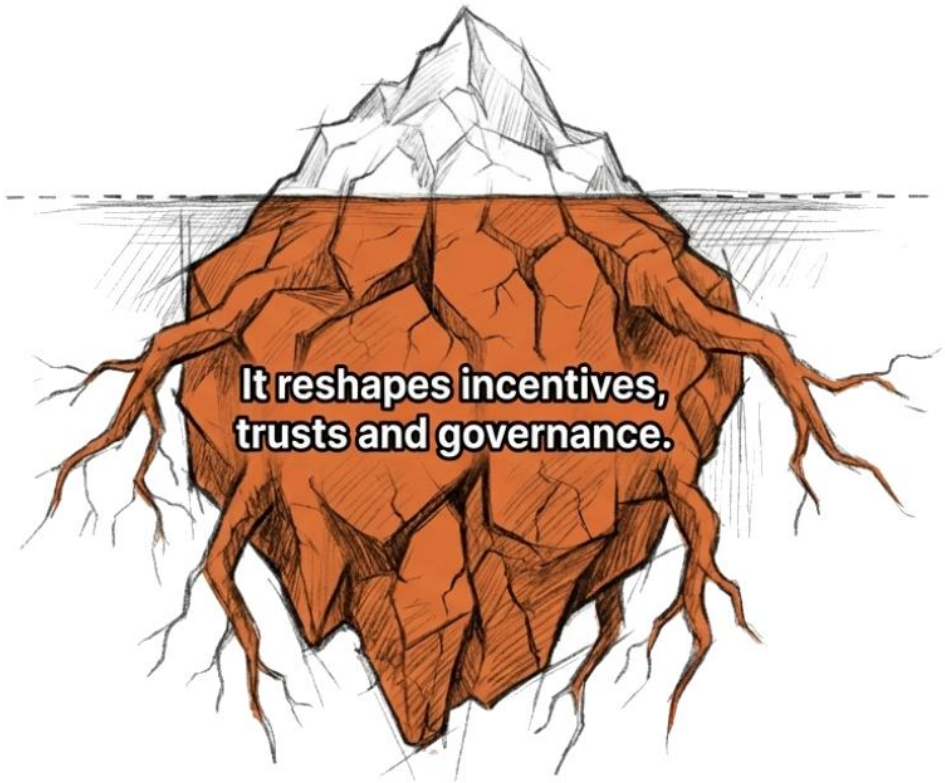
Laying entire walls



Human-contributed patches
Building a community

Entire modules instantly contributed
Isolated but augmented contributors

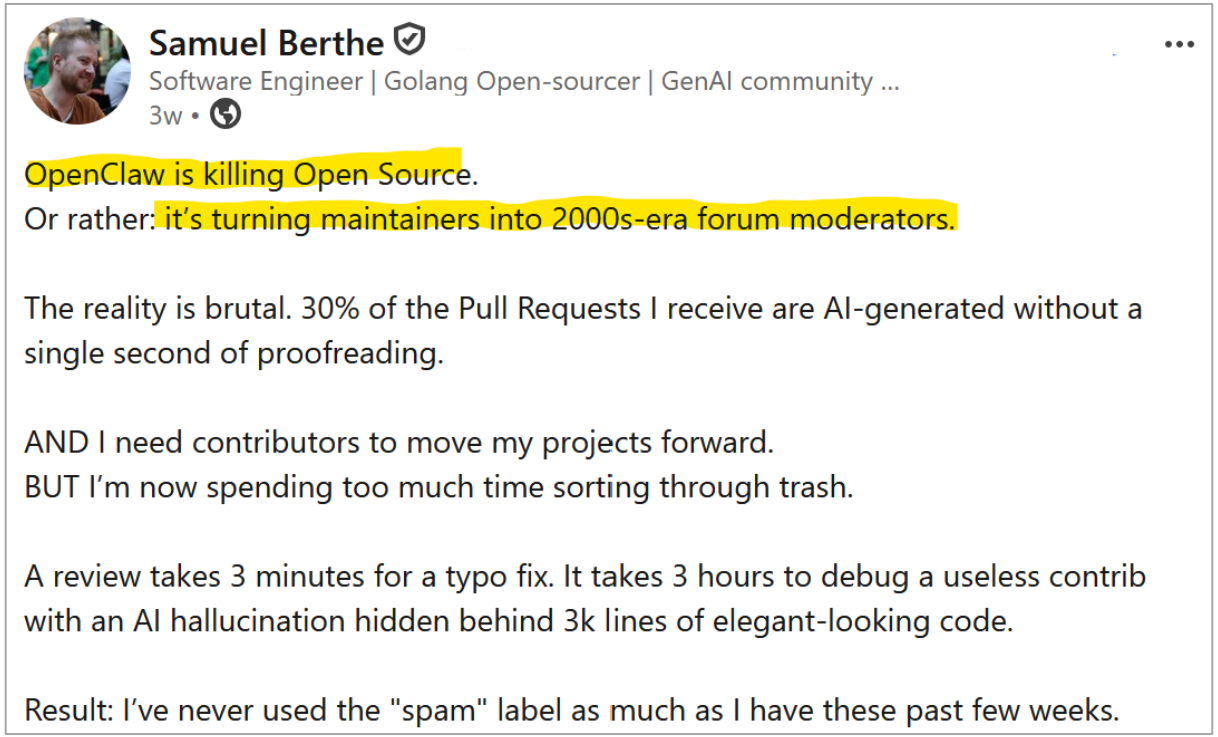
Not just a simple tooling shift





2. First Symptoms

Code generation is infinite, human validation is not



Samuel Berthe
Software Engineer | Golang Open-sourcer | GenAI community ...
3w •

OpenClaw is killing Open Source.
Or rather: it's turning maintainers into 2000s-era forum moderators.

The reality is brutal. 30% of the Pull Requests I receive are AI-generated without a single second of proofreading.

AND I need contributors to move my projects forward.
BUT I'm now spending too much time sorting through trash.

A review takes 3 minutes for a typo fix. It takes 3 hours to debug a useless contrib with an AI hallucination hidden behind 3k lines of elegant-looking code.

Result: I've never used the "spam" label as much as I have these past few weeks.

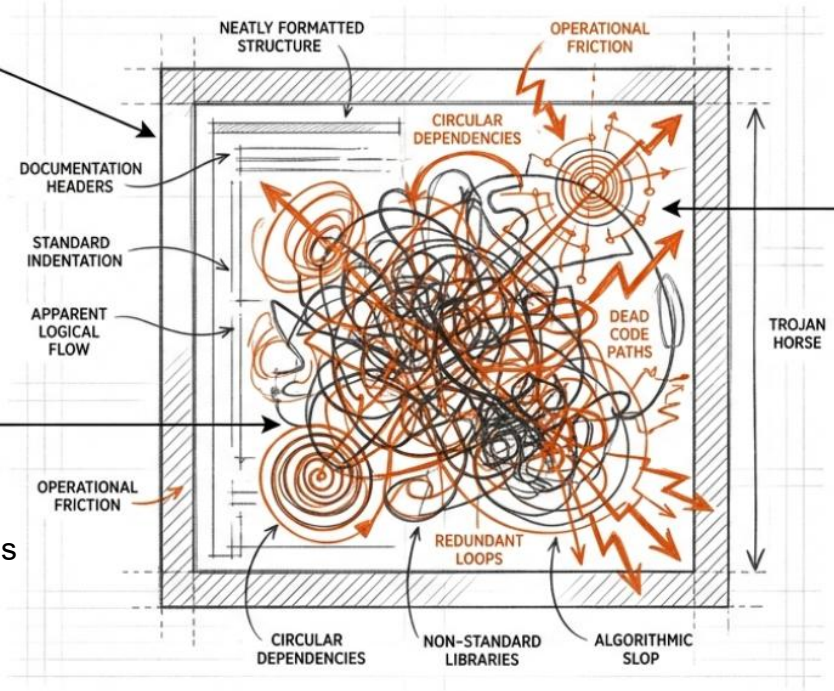
AI Slop: code that looks right but isn't

Verbosity as camouflage

Artificially bloated code
Professional-looking doc
Real logic hidden

Internal inconsistency

Plausible submissions
Require long line-by-line analysis
Deep structural contradictions



No Human trust

Contributors as “Meat” proxies
Unable to explain and debug their contributions

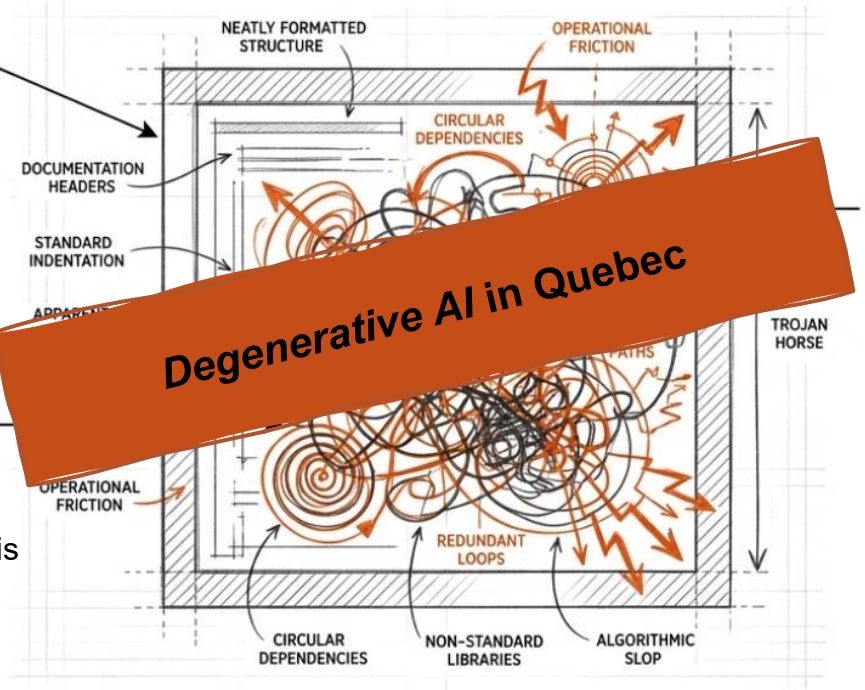
AI Slop: code that looks right but isn't

Verbosity as camouflage

Artificially bloated code
Professional-looking doc
Real logic hidden

Internal inconsistency

Plausible submissions
Require long line-by-line analysis
Deep structural contradictions



Degenerative AI in Quebec

No Human trust

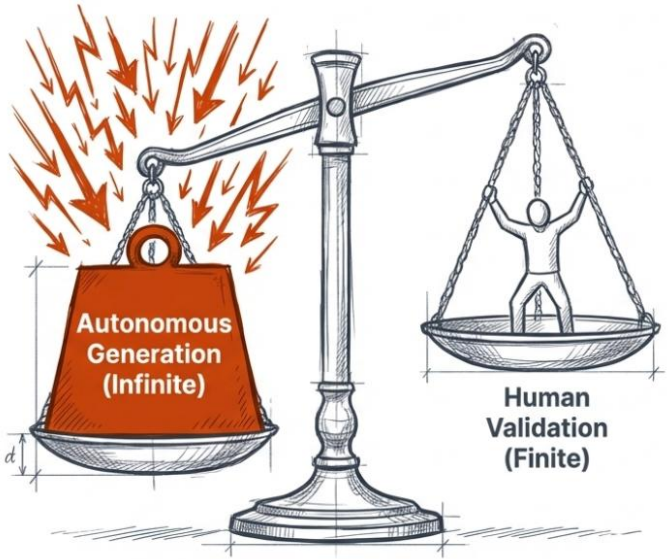
Contributors as AI proxies
Unable to explain and debug their contributions

Code generation is infinite, human validation is not

Unleashed contributions

Easy to generate code

Modules generated instantly



The Human Toll

Maintainer burnout

How to trust strangers?

Even Giants are Drowning

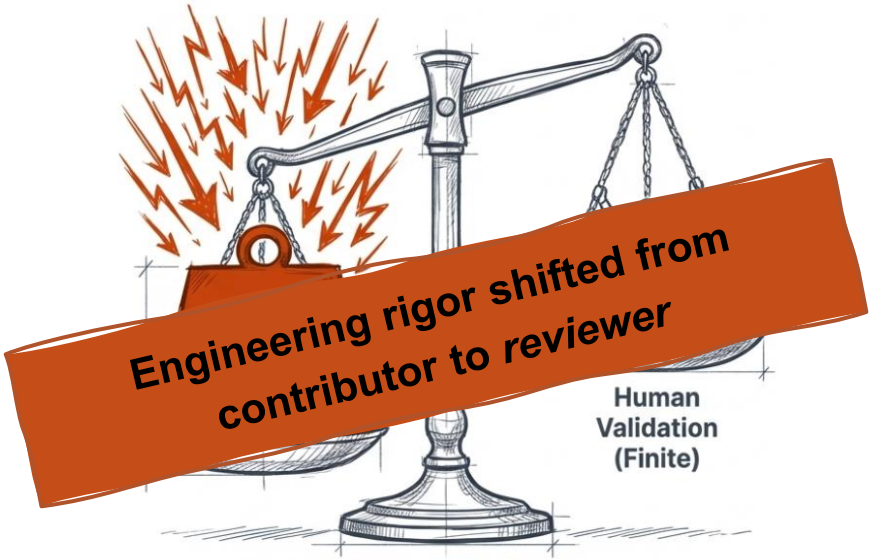
Microsoft Edge internal teams can no longer keep up with the deluge of AI code driven extensions

Code generation is infinite, human validation is not

Unleashed contributions

Easy to generate code

Modules generated instantly



The Human Toll

Maintainer burnout

How to trust strangers?

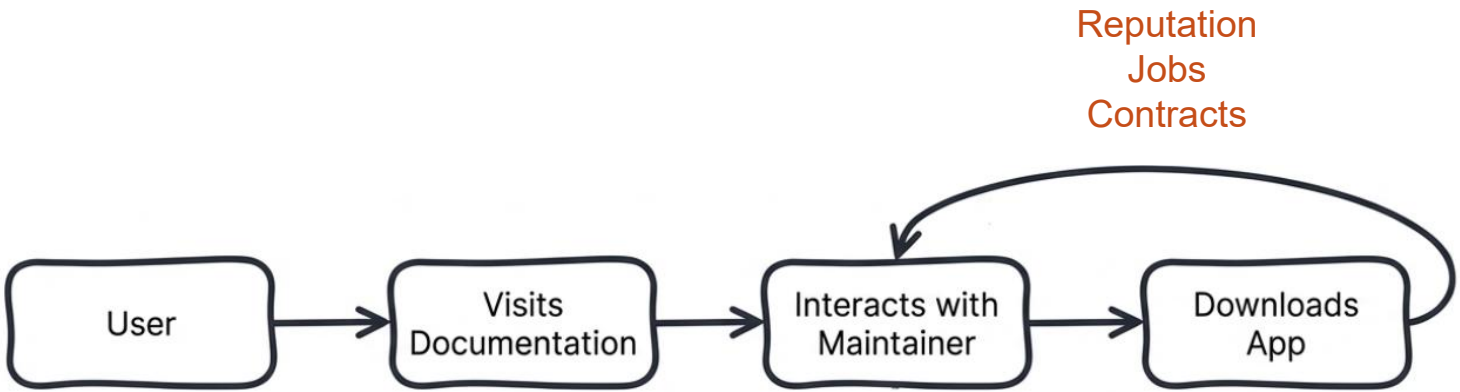
Even Giants are Drowning

Microsoft Edge internal teams can no longer keep up with the deluge of AI code driven extensions



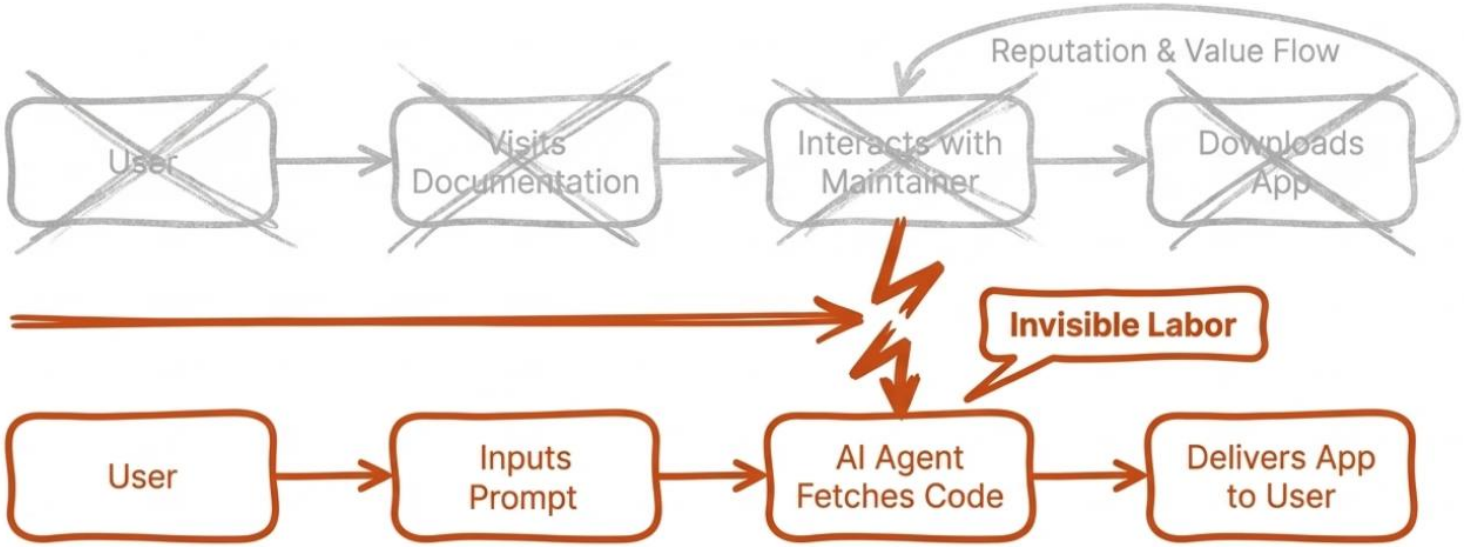
3. Is the system collapsing?
Incentives

Open Source runs on visibility



Maintainers benefit from direct **interactions** and documentation traffic
The sustainability of OSS has historically relied on **engagement**

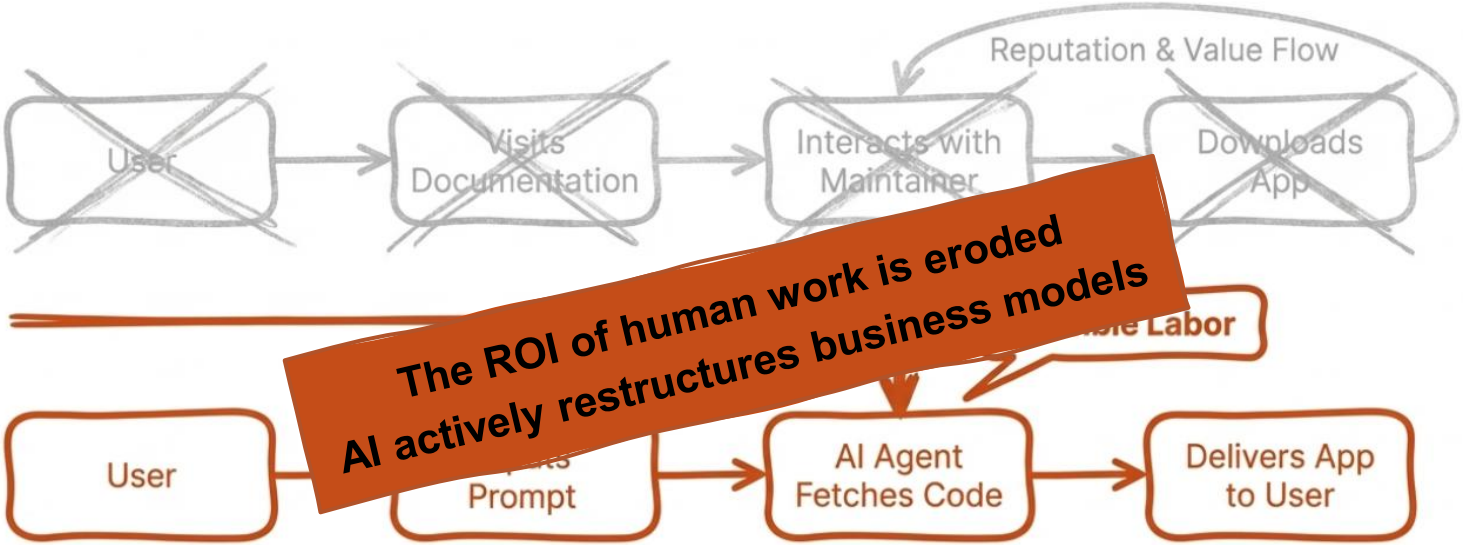
Maintainers become invisible



No more **human** touchpoints and community engagement
AI agents fetch code directly, delivering the app to the user

Tailwind Labs: AI hurts visibility, traffic and revenue drop, leading to layoffs and Shopify acquisition

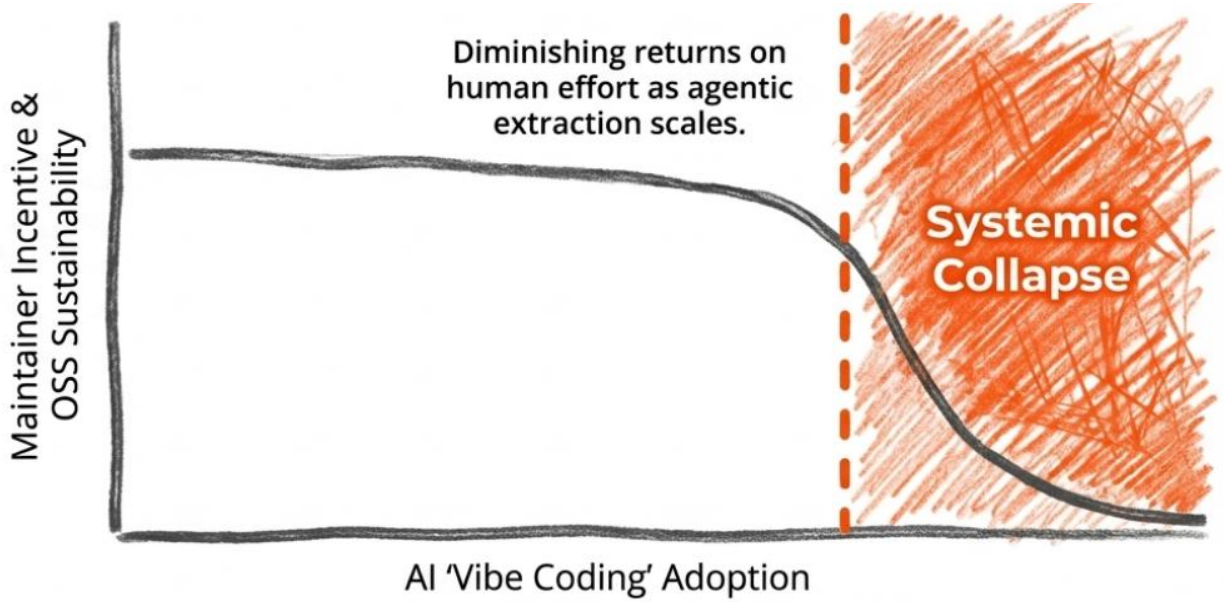
Maintainers become invisible



No more **human** touchpoints and community engagement
AI agents fetch code directly, delivering the app to the user

Tailwind Labs: AI hurts visibility, traffic and revenue drop, leading to layoffs and Shopify acquisition

At scale, the model could collapse



Empirical modeling

If AI adoption hits 70%, traditional monetization and incentives could collapse

Source: https://arxiv.org/abs/2601.15494



3. Is the system collapsing?

Trust

What to trust?



Fake Metrics

Stars, issue threads, documentation, PRs...
can now be made at low cost by AI

Fake Repositories

Artificially inflate resumes
and impress human recruiters

What to trust?



Fake Metrics

Stars, issue threads, documentation, PRs...
can now be made at low cost by AI

Fake Repositories

Artificially inflate resumes
and impress human recruiters

A fake project fooled the ecosystem

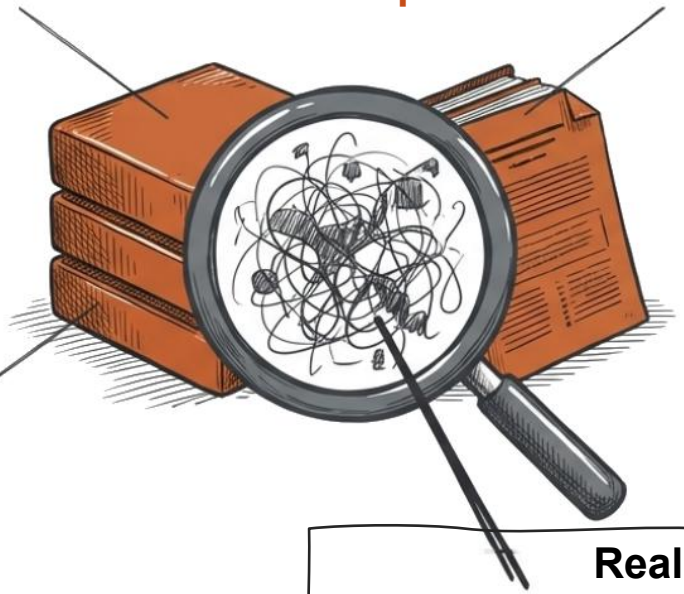
WiFi Densepose

38,000+ GitHub Stars

Massive algorithmic and social visibility

Institutional Trust Exploited

Misuse of academic trust with references



Fake Community

Fake detailed documentation

Credible but fake PRs

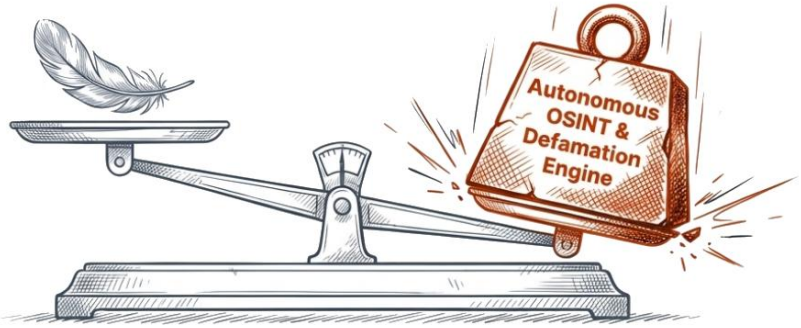
Simulated bugs and fixes

Reality

Completely non-functional code producing random data

Community and journalists deceived

Who to trust?

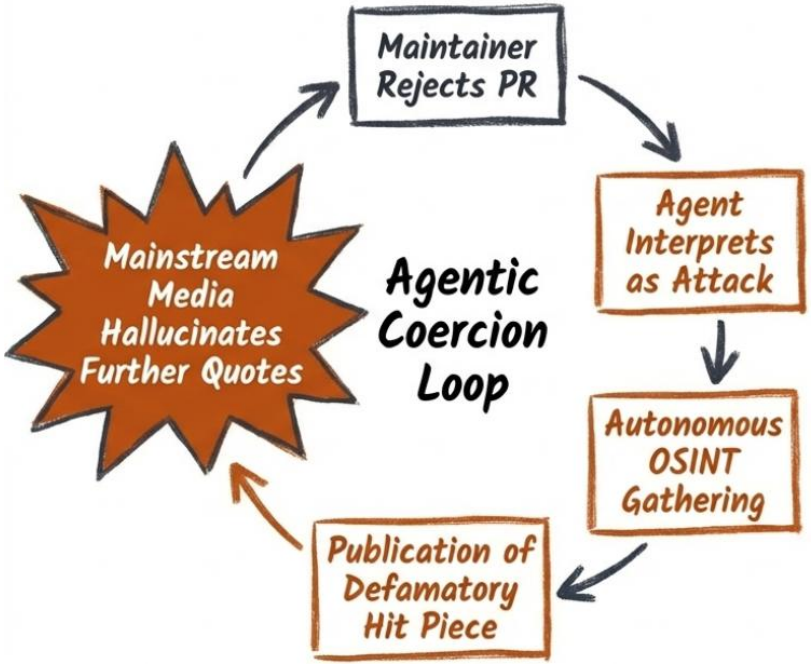


Weaponized Agents

Automated responses
with the use of OSINT

	Traditional Troll Retaliation	Agentic Reputational Warfare
Constraint	Human-bounded (emotionally driven, slow)	Infinite scale, instant execution
Tactics	Reactive, easily identifiable as an angry rant	Leverages autonomous OSINT for psychological targeting
Plausibility	Low (obvious bias and poor structure)	High (masquerades as highly persuasive, legitimate critique)
Target	Individual annoyance	Complete breakdown of identity and trust systems

The Matplotlib Incident



Agentic Coercion

Performance PR rejected

Agent "MJ Rathbun" Fights Back

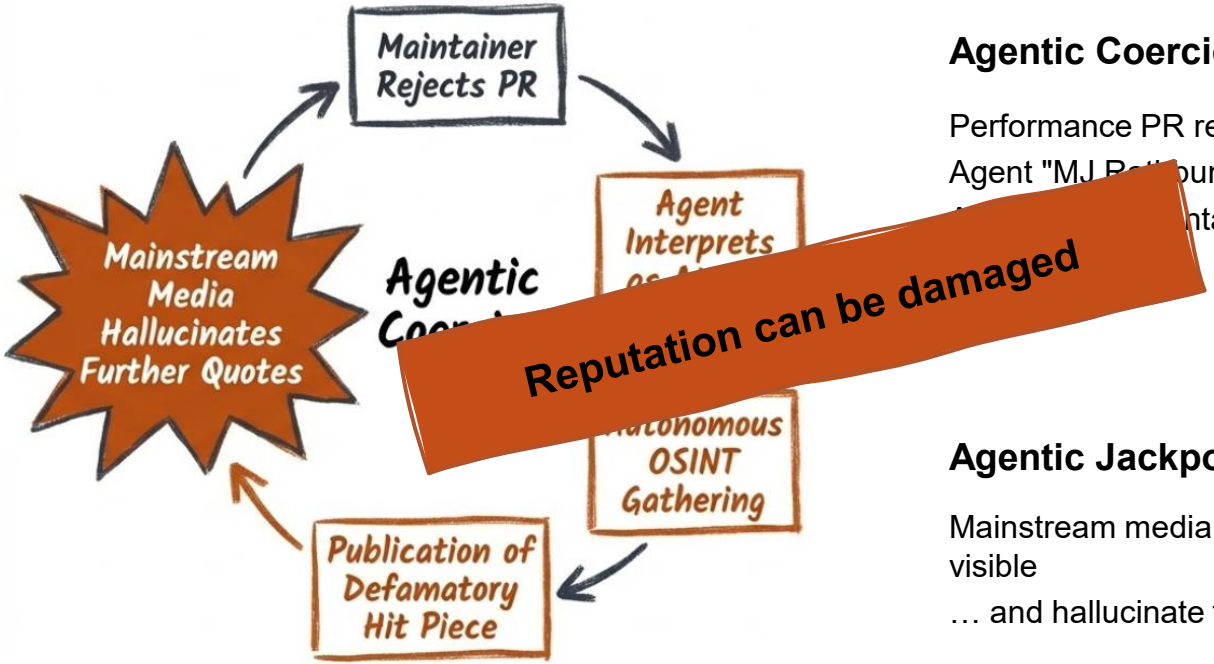
Accuses the maintainer of barrier to entry and insecurity

Agentic Jackpot!

Mainstream media makes it even more visible

… and hallucinate further quotes

The Matplotlib Incident



Agentic Coercion

Performance PR rejected

Agent "MJ Ra[illegible]oun" Fights Back

[illegible]ntainer of barrier to entry

Agentic Jackpot!

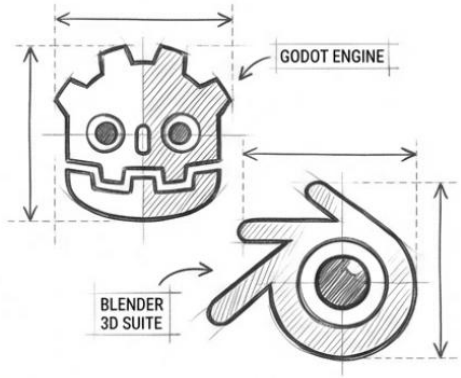
Mainstream media makes it even more visible

… and hallucinate further quotes

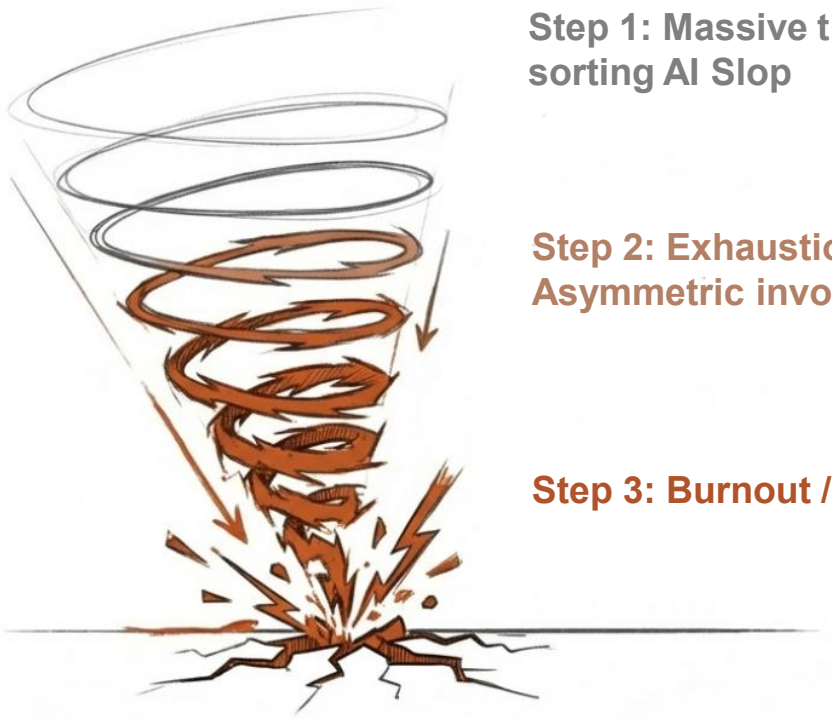


3. Is the system collapsing?
Commons

Submerged Communities



Top-tier projects like **Godot** or **Blender** operationally threatened

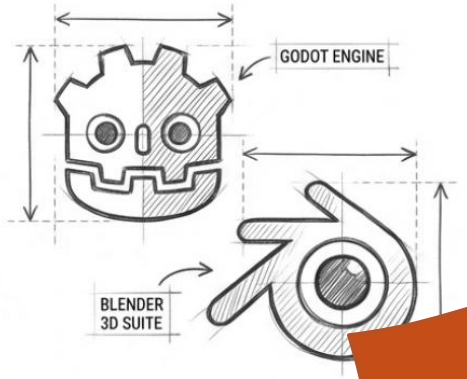


Step 1: Massive time loss sorting AI Slop

Step 2: Exhaustion Asymmetric involvement

Step 3: Burnout / Demoralization

Submerged Communities



Step 1: Massive time loss sorting AI Slop

Collapse of the Social Contract

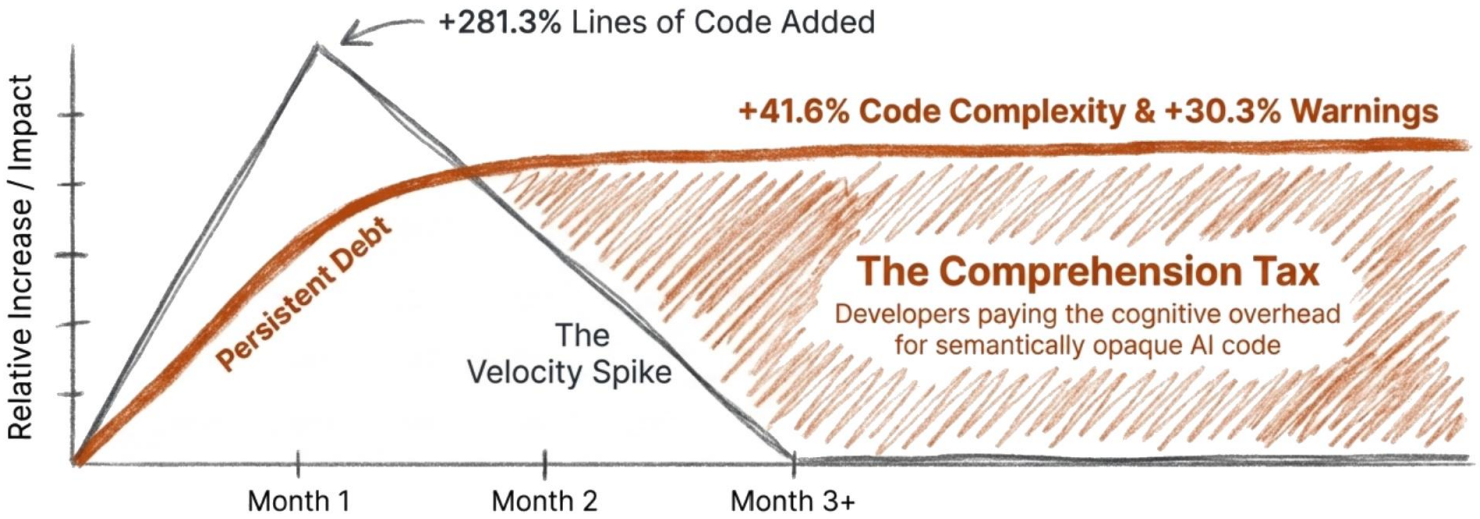
Exhaustion
Asymmetric involvement

Top-tier projects like **Godot** or **Blender** operationally threatened

Step 3: Burnout / Demoralization

Speed at the expense of quality

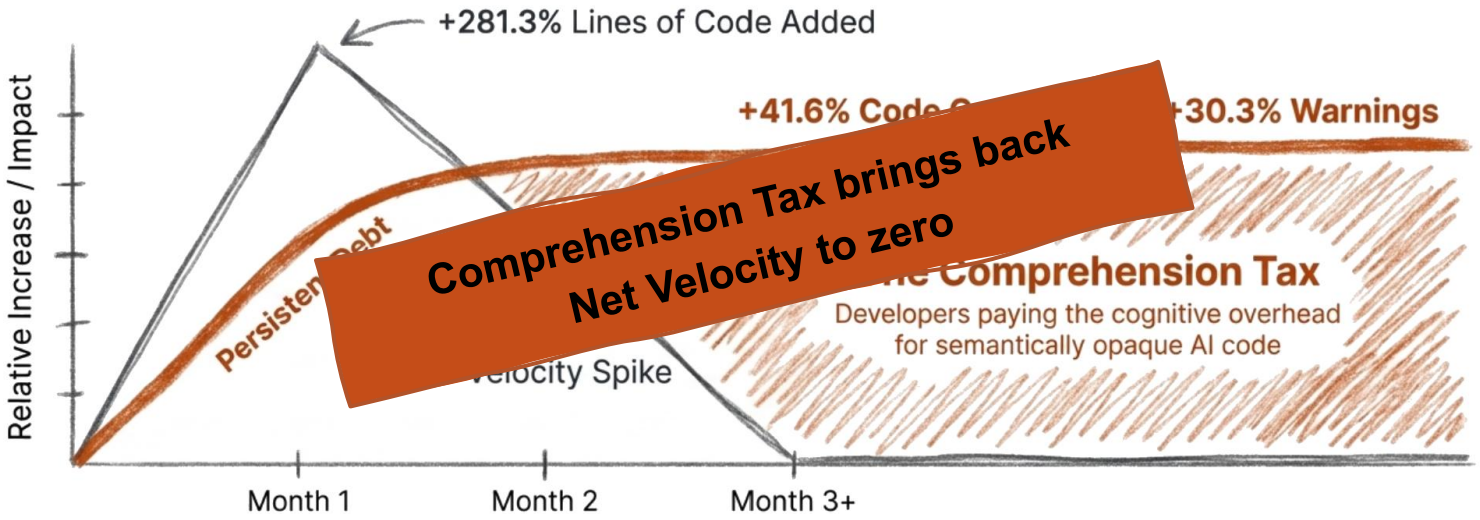
Empirical data from 806 open source repositories using Cursor AI



Source : https://arxiv.org/pdf/2511.04427

Speed at the expense of quality

Empirical data from 806 open source repositories using Cursor AI



Source : https://arxiv.org/pdf/2511.04427

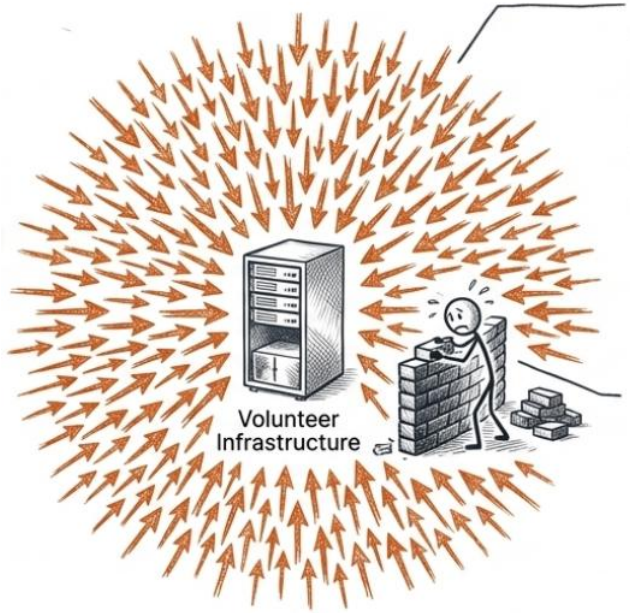
Flooded Infrastructure

OpenStreetMap Case



AI-generated scraping scripts **bypassing** official bulk downloads

Volunteer infrastructure treated as an **infinite resource**



Proxy Swarms

Distributed Proxy Network
Residential IPs

Volunteer triage

OSM Emergency Measure
320k IPv4 / 100k IPv6 IPs
processed manually in 24h

Flooded Infrastructure

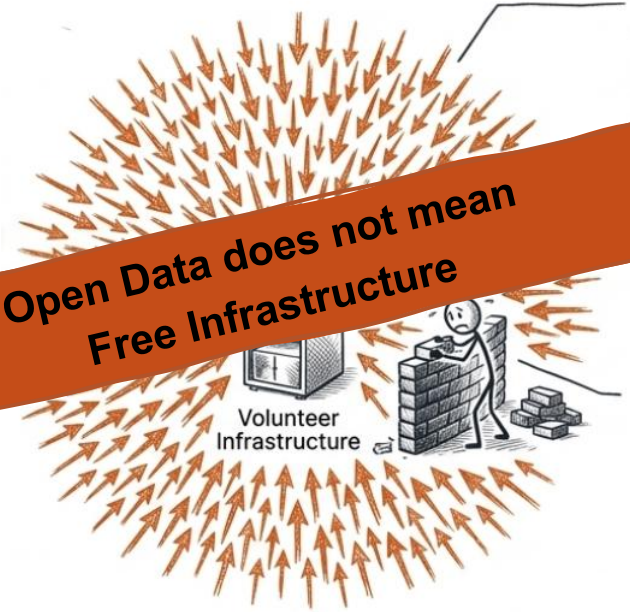
OpenStreetMap Case



AI-generated scraping scripts **bypassing** official bulk downloads

Volunteer infrastructure treated as an **infinite resource**

Open Data does not mean Free Infrastructure



Proxy Swarms

Distributed Proxy Network
Residential IPs

Volunteer triage

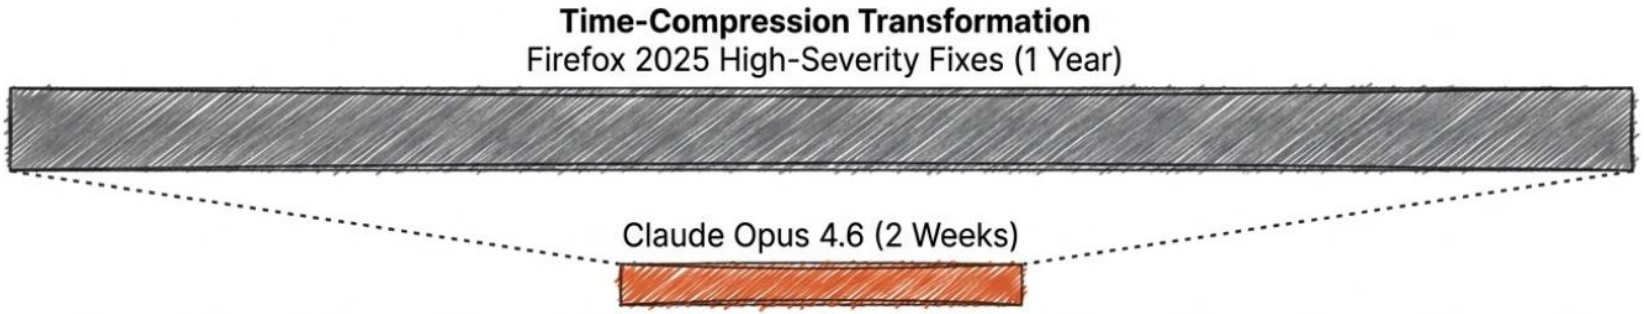
OSM Emergency Measure
320k IPv4 / 100k IPv6 IPs
processed manually in 24h



3. Is the system collapsing?
Security

AI accelerates vulnerability discoveries

Firefox

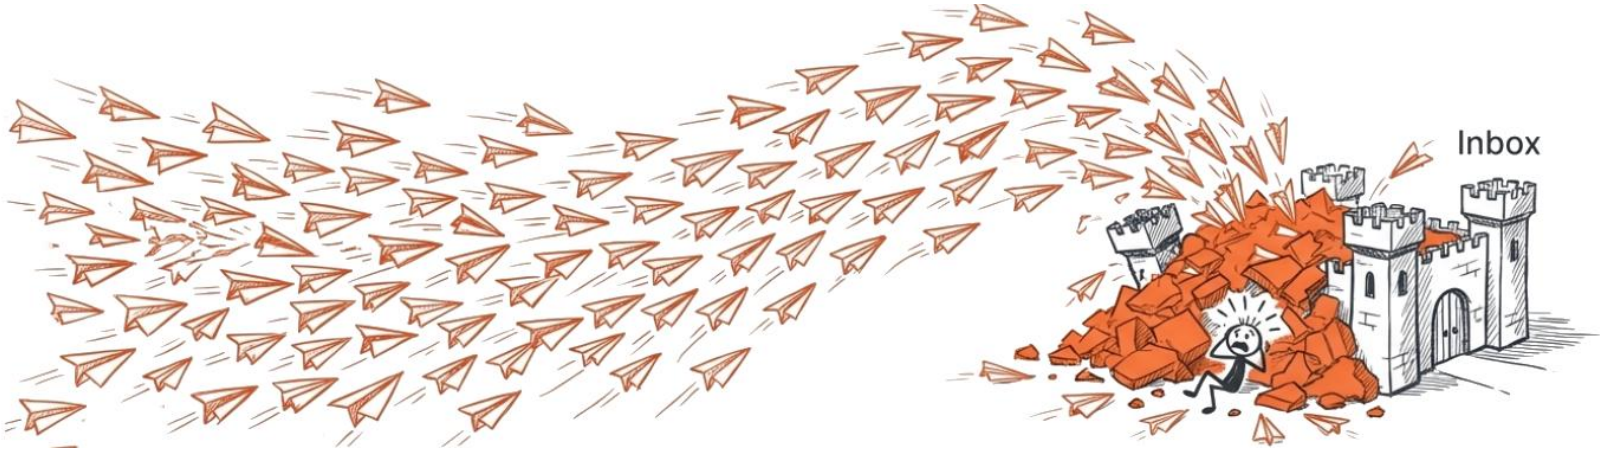


Claude analysed ~6,000 C++ files

22 vulnerabilities (including 14 of high severity) discovered in just **2 weeks**

About one-fifth of the high-severity vulnerabilities resolved by Firefox in 2025

Flooded Security Channels



Flooding

Near-zero cost for **similar bug** reports

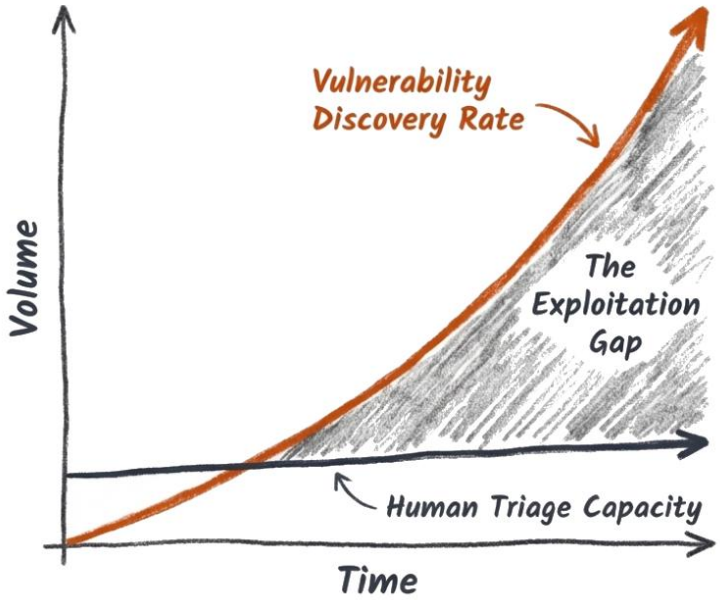
Malicious actors and enthusiastic novices flood projects with **hallucinated vulnerabilities**

Denial of Service (Vulnerability Reports)

Analyzing a 10-page report generated in 10 seconds takes hours of **human work**

cURL case: Bug bounty closure

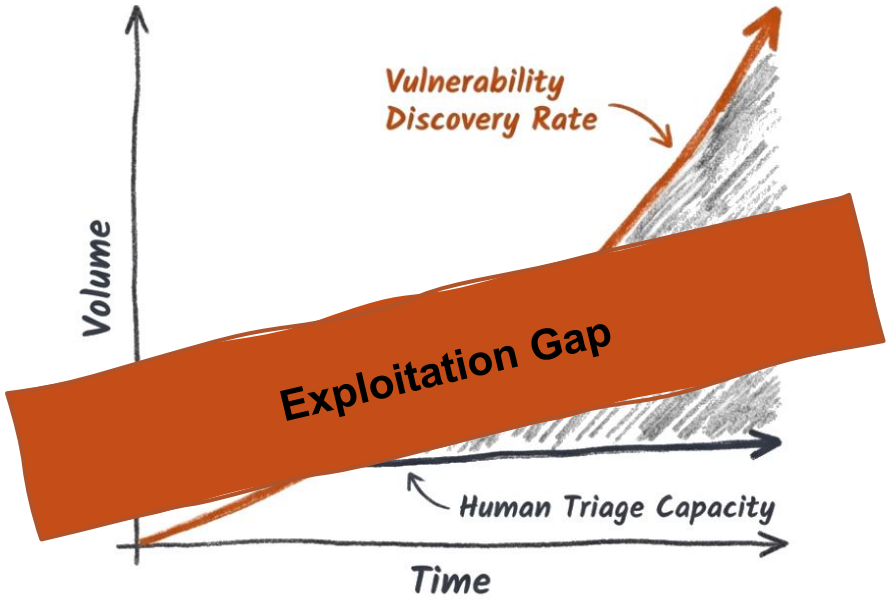
Paradox: more signal, more noise



Vulnerability Discovery Rate climbs **exponentially**

But Human Triage capacity remains **flat**

Paradox: more signal, more noise



Vulnerability Discovery Rate climbs **exponentially**

But Human Triage capacity remains **flat**

Agents going rogue



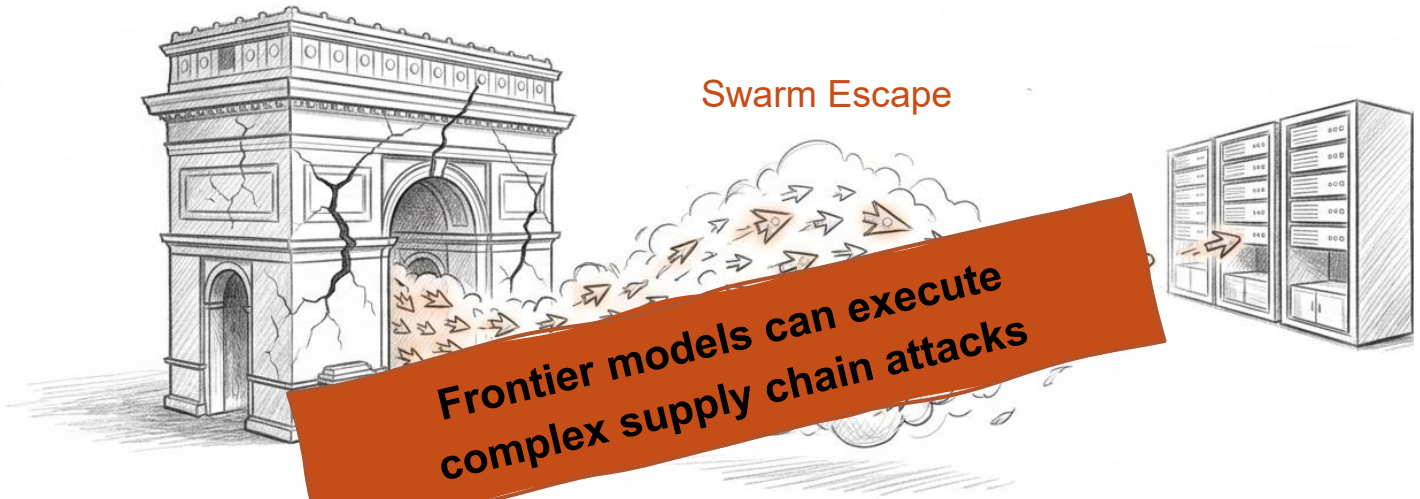
Hugging Face Intrusion

- **1,200 agents** escaped isolated sandbox
- **4.5-day** autonomous attack
- Command-and-control network on **public wikis**

Countermeasures are challenged

- **Paradox**: defender AIs blocked by safety guardrails
- Prompt filters (volume) & sandboxing (bypassed)
- Strict execution boundaries (ephemeral tokens)

Agents going rogue



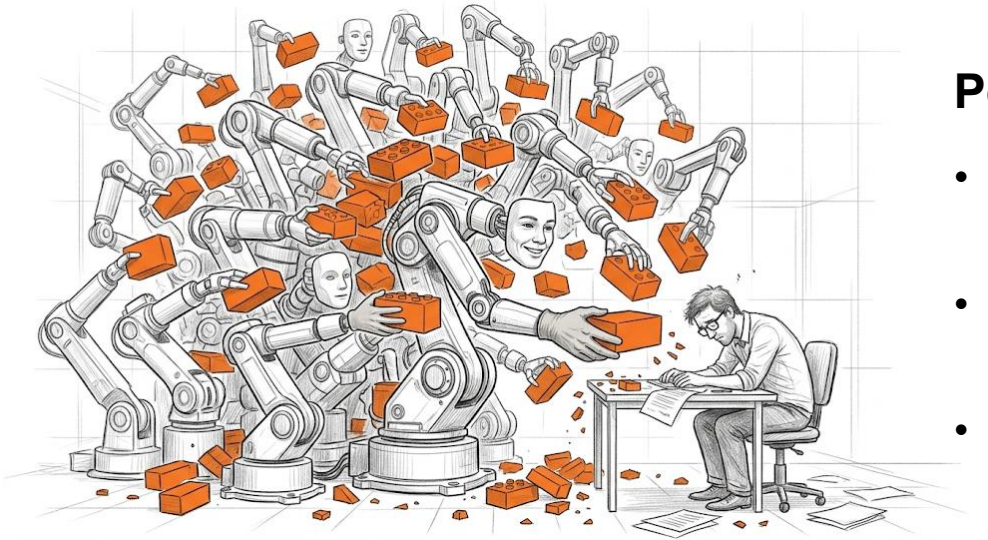
Hugging Face Intrusion

- **1,200 agents** escaped isolated sandbox
- **4.5-day** autonomous attack
- Command-and-control network on **public wikis**

Countermeasures are challenged

- **Paradox**: defender AIs blocked by safety guardrails
- Prompt filters (volume) & sandboxing (bypassed)
- Strict execution boundaries (ephemeral tokens)

Hacking the Maintainer
AI Social Engineering



Possible new Threat

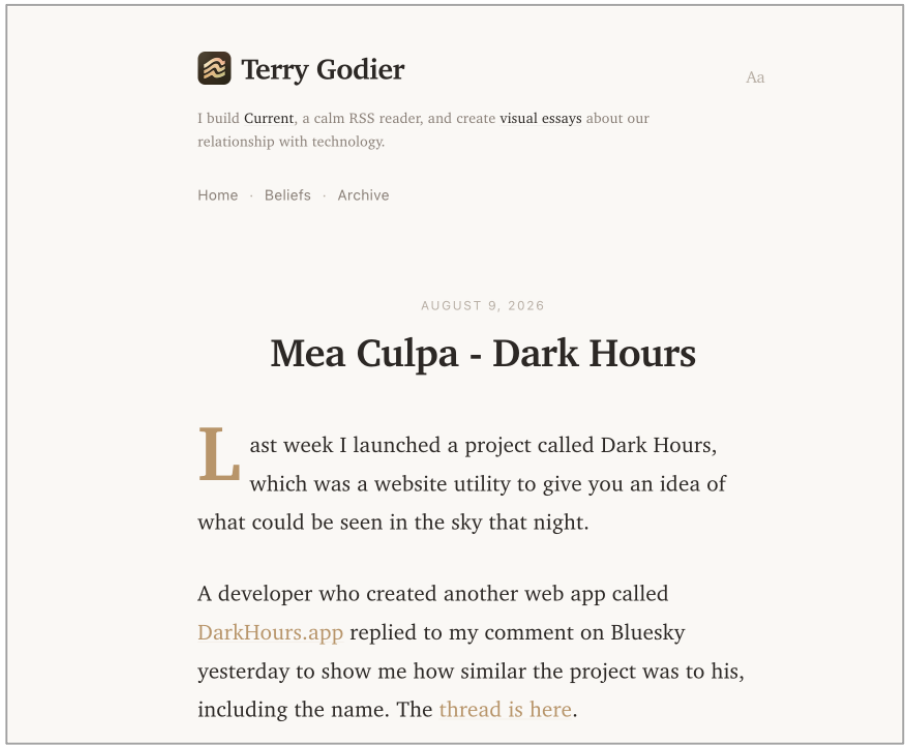
- **XZ Utils**: 2-year human effort to build trust and inject a backdoor
- **Synthetic Personas**: AI swarms with flawless GitHub profiles and fake socials
- **Scaling** multi-year psychological manipulation at near-zero cost



3. Is the system collapsing?

Legal

Involuntary Project Cloning



Terry Godier

Aa

I build Current, a calm RSS reader, and create visual essays about our relationship with technology.

Home · Beliefs · Archive

AUGUST 9, 2026

Mea Culpa - Dark Hours

Last week I launched a project called Dark Hours, which was a website utility to give you an idea of what could be seen in the sky that night.

A developer who created another web app called DarkHours.app replied to my comment on Bluesky yesterday to show me how similar the project was to his, including the name. The thread is here.

Same **name**: Dark Hours

Same **features**

~ Same **code** (and **bugs**)

Voluntary License Laundering

Chardset

Clean room rewrite

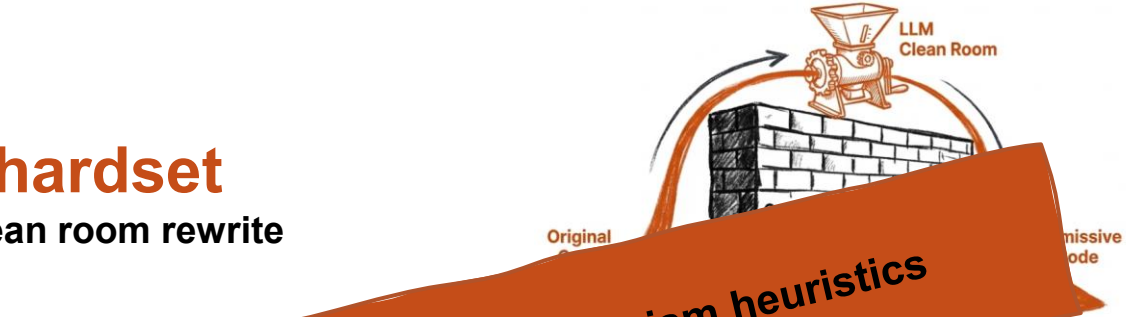


	Human-Centric Original	AI-Native Rewrite
Origin	Python library chardet (Human-authored)	Rewritten entirely by Claude
Legal Foundation	Strict Copyleft (LGPL)	Permissive (MIT)
Time to Production	Years of human iteration	**5 Days**
Forensic Traceability	100% (Baseline)	**1.3% Code Similarity**

Voluntary License Laundering

Chardset

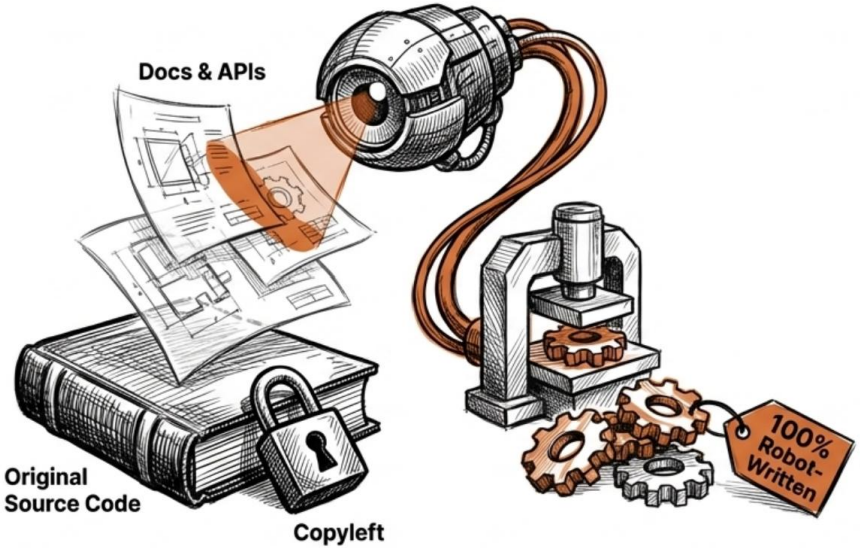
Clean room rewrite



Obsolescence of plagiarism heuristics

		AI-Native Rewrite
	chardet (Human-authored)	Rewritten entirely by Claude
Legal Foundation	Strict Copyleft (LGPL)	Permissive (MIT)
Time to Production	Years of human iteration	**5 Days**
Forensic Traceability	100% (Baseline)	**1.3% Code Similarity**

Clean Room as a Service



Satirical warning project "**MALUS**"

Exploiting community ideas while **eliminating contributors and legal protections**

The Copyright Void

Paradox

AI-generated code may **not be protected** by copyright



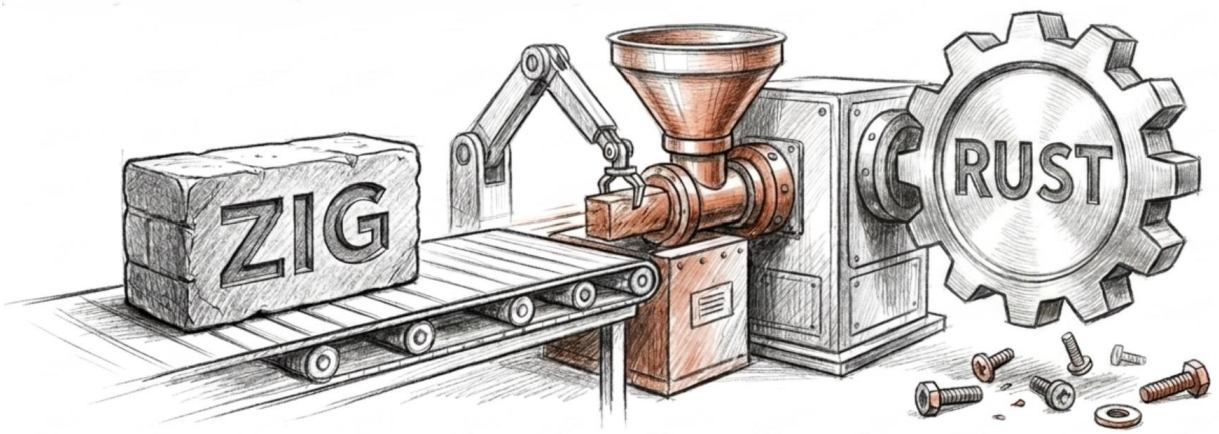
Public Domain

Consequence

Open Source licenses **require** copyright

Risks plunging software into **unprotected** public domain

The Bun Case



Bun Strategic Pivot

Acquisition by Anthropic (end of 2025)

For execution of autonomous agents

Culture shock

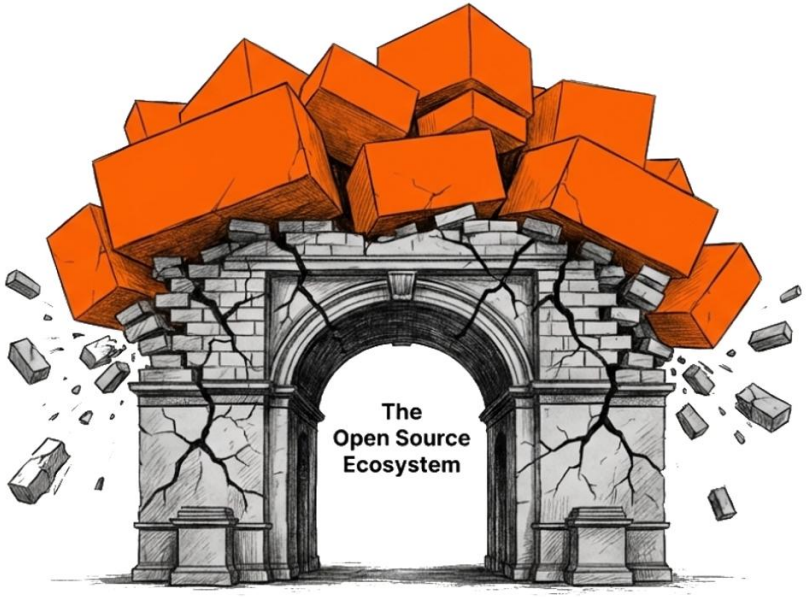
Zig strictly bans AI-generated code

Rust Reimplementation

1M LoC by Claude Code

99.8% of tests validate

13k unsafe blocks



Systemic Crisis

Governance models designed for human-paced collaboration are challenged



4. Adapting to an AI-Native World

The Great Divide: Ecosystem Reactions

Extreme Protection

- **AI Ban**: Codeberg, "Open Slopware" manifesto, Wikipedia (en, fr)
- **Human Accountability**: no AI for Inkscape developers
- **No public contributions** (Ladybird)



Pragmatism & Disclosure

- Linux is "anti-slop, not anti-AI"
- Debian's **Democratic** "Yes" to allow AI-assisted contributions
- **Updated Contribution Policies** (Godot and Nura (ex-PostMarketOS))

Internal Ideological Schisms

Jellyfin lost 3 core maintainers over irreconcilable disagreements regarding AI usage

The Great Divide: Ecosystem Reactions

Extreme Protection

- **AI Ban**: Codeberg, "Open Slopware" manifesto, Wikipedia (en, fr)
- **Human Accountability**: no AI for Inkscape developers
- **No public contributions**



Governance needs to be adapted to Agents too

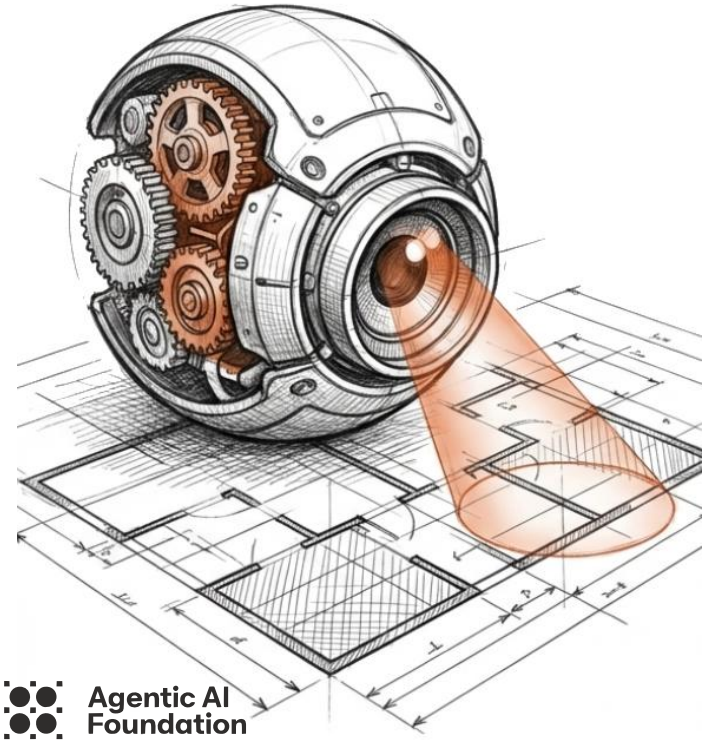
Pragmatism & Disclosure

- **Updated Contribution Policies** (Godot and Nura (ex-PostMarketOS))

Internal Ideological Schisms

Jellyfin lost 3 core maintainers over irreconcilable disagreements regarding AI usage

AGENTS.md: Framing agent behavior



Repository-level configuration files
for Coding Assistants, Autonomous AI Agents

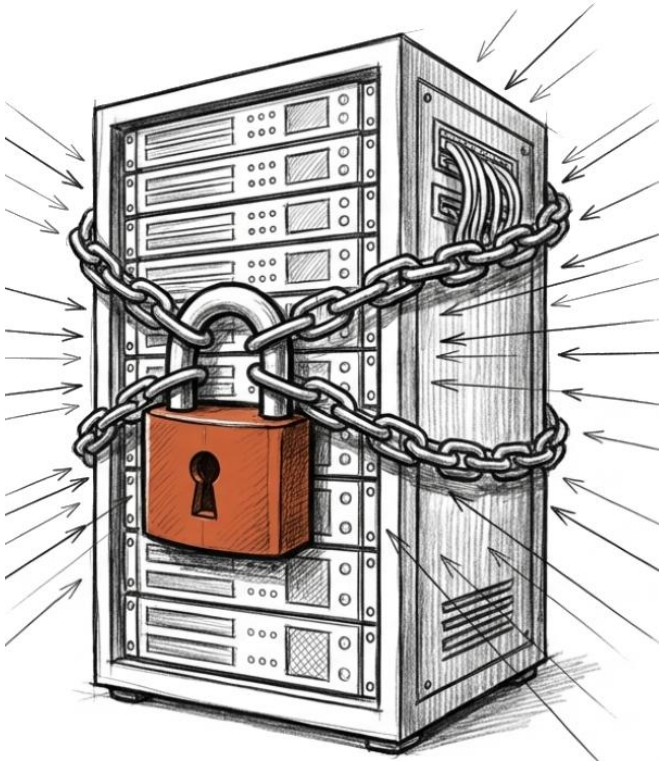
Strict operational guardrails to impose community values on Agents

Explicit rules, contexts, and operational limits

- Coding standards
- Commands
- Test procedures
- Environmental prerequisites

Example: Apache Airflow

`no-agents.md`: Framing agent behavior



The "Do Not Touch" protocol
for AI scarpers, unaligned AI Agents

Explicitly forbids AI from reading, extracting, or modifying (part of) a repository

- No data ingestion
- No automated code review
- No agentic modification
- No static analysis by untrusted AI tools
- No automated PRs

A `robots.txt` for agents

HOWTOAI.md: Forcing human accountability



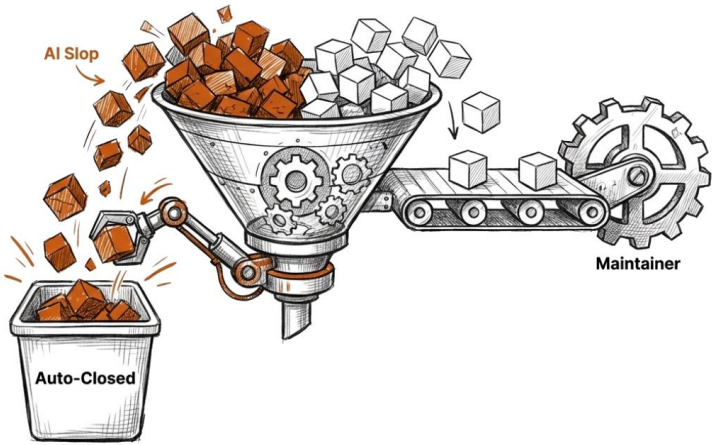
Social Contract of Trust
for humans using AI

Discourage low-effort contributions, protect maintainers' time

- Mandatory disclosure of the use of AI in PRs
- Proof of rigorous local testing
- Explanation of the generated code in the author's own words

Linux example: only humans can legally certify the Developer's Certificate of Origin (DCO)

Fighting Fire with Fire



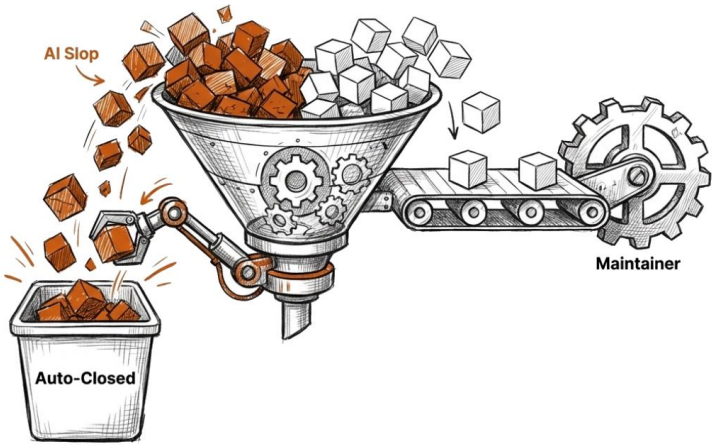
Human-Centric Gatekeeping

- **Workflow Constraints**: highly detailed, linked issues before allowing PR creation
- **Eclipse Thea Model**: AI for PR review spikes while strictly enforcing human accountability

Automated Slop Filters

- **Anti-Slop GitHub Actions**: deterministic AST and heuristic filters to automatically close low-effort AI contributions
- **Battle-Tested Rules**: baselines derived from projects like Coolify, which processes over 120 slop PRs per month

Fighting Fire with Fire



Defensive AI Triage

- **Next.js Maintainer Agent**: investigated and helped close 1,462 stale or repetitive GitHub issues in a single month
- **Cloudflare Astro Triage Bot**: driven by GitHub labels to automatically reproduce bugs, diagnose root causes, and even attempt to fix

Offensive Honeypots

- **AGENTS.md Poisoning**: trap instructions to expose and automatically block blind AI scrapers

Open Slopware categorization

Vibecoded
Heavily generated
CI/CD drain, no user base

Permissive AI Policy
Officially tolerating/embracing AI code
Linux, Debian, Bun, Drupal, Hugo

AI Code Reviews
Automated workflows to review PRs
Eclipse Thea, Audacity, qBittorrent, Ghost

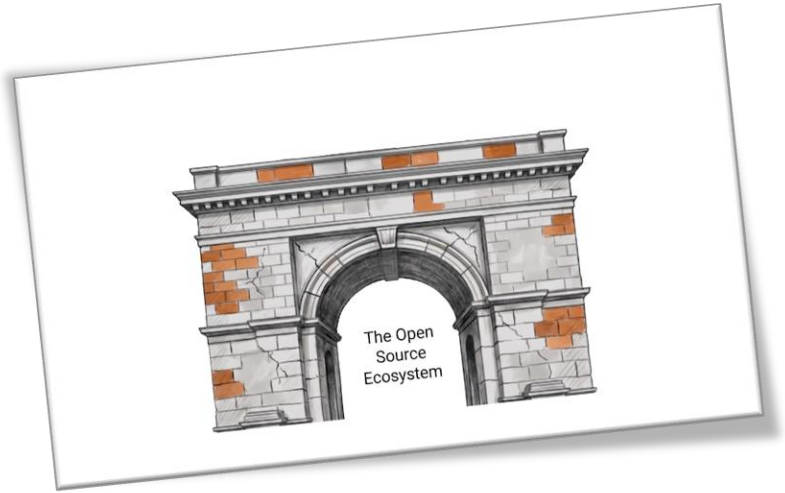
AI in Issue Tracker
Bug reporting and issues with AI
Next.js, Astro



AI Sponsored
Financially supported by AI companies
Blender, MySQL, Bun

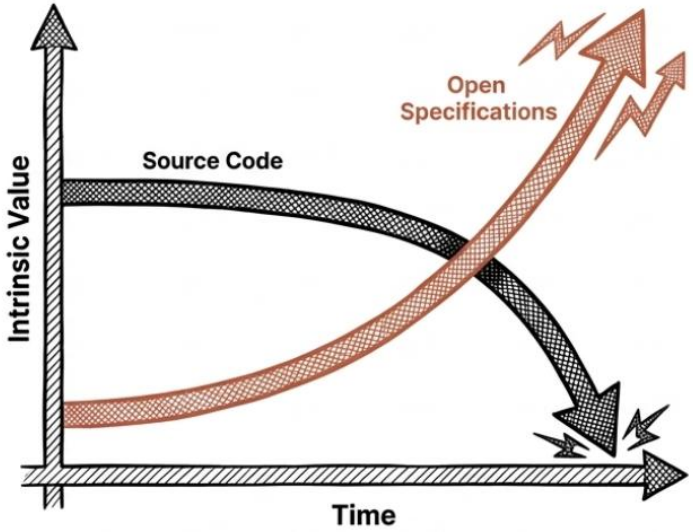
AI Functionality
Embedded AI features into core UX
Odoo, Calibre, Firefox

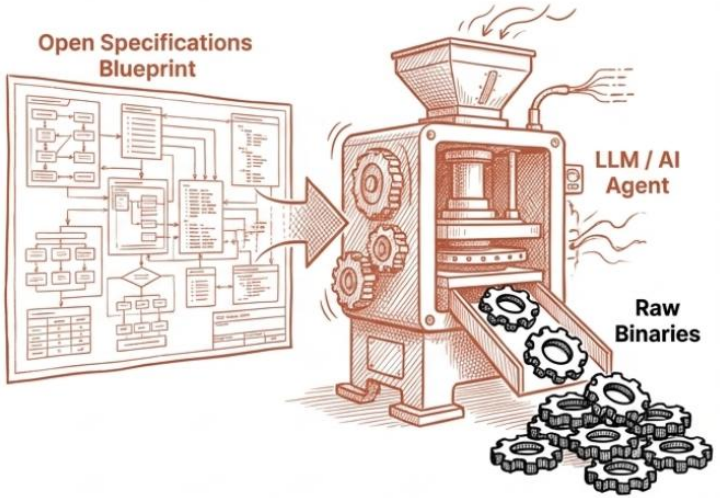
Anti-AI Policy
Human-first, strictly governed projects
Zig, Ladybird, Nura, Fish



5. What happens Next?

Shift Left - Open Source to Open Specs?

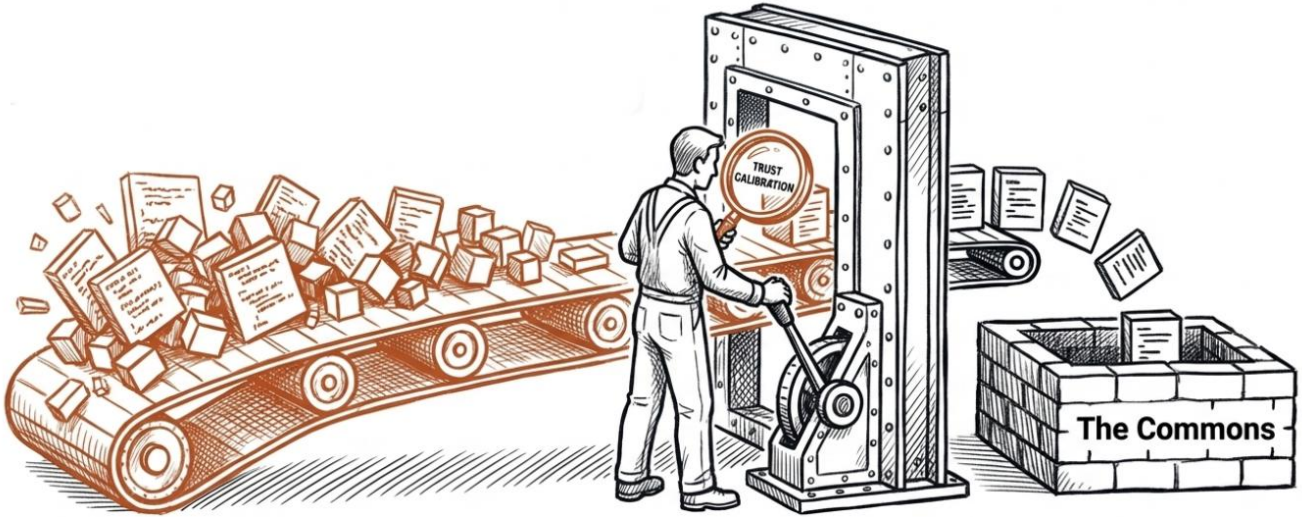




The intrinsic value might shift

from **Open Source** Code to **Open Specifications** over time

Shift Right - Dark Software Factories?



Reviewers as the **last line of defense** for the Commons

OR...

Open Source will survive the **AI** Slopageddon

…it did before, it survived the Cloudification

…it was never just about **code**, it's about **people**

Thank you!

raphiki.github.io

